

***** KUBERNETES_INSTALLATION*****

Install Kubernetes Cluster on CentOS 7 with kubeadm and CRI-O

create 3 vm of centos image

Server Type	Server Hostname	Specs
Master	master-node	4GB Ram, 2vcpus
Worker	node1	2GB Ram, 1vcpus
Worker	node2	2GB Ram, 1vcpus

edit host file and write all ip-address of all machine on all system

vi /etc/hosts

```
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1         localhost localhost.localdomain localhost6 localhost6.localdomain6
192.168.174.166 master
192.168.174.167 node1
192.168.174.168 node2
```

make sure all 3 pc ping to each other

create keygen on master node

Ssh-keygen

copy keygen to all client node

Ssh-copy-id

Update all system

yum update

disable selinux on all system

setenforce 0

'or'

vi /etc/selinux/config

SELINUX=disabled

```
# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
#     enforcing - SELinux security policy is enforced.
#     permissive - SELinux prints warnings instead of enforcing.
#     disabled - No SELinux policy is loaded.
SELINUX=disabled
# SELINUXTYPE= can take one of three values:
#     targeted - Targeted processes are protected,
#     minimum - Modification of targeted policy. Only selected processes are protected.
#     mls - Multi Level Security protection.
SELINUXTYPE=targeted
```

firewall disable on all system

```
systemctl disable firewalld
systemctl stop firewalld
```

restat all system on all machine

```
init 6
```

```

[root@node2 ~]# modprobe overlay
[root@node2 ~]# modprobe br_netfilter
[root@node2 ~]# tee /etc/sysctl.d/kubernetes.conf<<EOF
> net.bridge.bridge-nf-call-ip6tables = 1
> net.bridge.bridge-nf-call-iptables = 1
> net.ipv4.ip_forward = 1
> EOF
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
[root@node2 ~]# sysctl --system
* Applying /usr/lib/sysctl.d/00-system.conf ...
net.bridge.bridge-nf-call-ip6tables = 0
net.bridge.bridge-nf-call-iptables = 0
net.bridge.bridge-nf-call-arptables = 0
* Applying /usr/lib/sysctl.d/10-default-yama-scope.conf ...
kernel.yama.ptrace_scope = 0
* Applying /usr/lib/sysctl.d/50-default.conf ...
kernel.sysrq = 16
kernel.core_uses_pid = 1
kernel.kptr_restrict = 1
net.ipv4.conf.default.rp_filter = 1
net.ipv4.conf.all.rp_filter = 1
net.ipv4.conf.default.accept_source_route = 0
net.ipv4.conf.all.accept_source_route = 0
net.ipv4.conf.default.promote_secondaries = 1
net.ipv4.conf.all.promote_secondaries = 1
fs.protected_hardlinks = 1
fs.protected_symlinks = 1
* Applying /usr/lib/sysctl.d/60-libvirt.conf ...
fs.aio-max-nr = 1048576
* Applying /etc/sysctl.d/99-sysctl.conf ...
* Applying /etc/sysctl.d/kubernetes.conf ...
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
* Applying /etc/sysctl.conf ...
[root@node2 ~]# █

```

add repo file on all machine

```
tee /etc/yum.repos.d/kubernetes.repo<<EOF
```

```
[kubernetes]
```

```
name=Kubernetes
```

```
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
```

```
enabled=1
```

```
gpgcheck=1
```

```
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
EOF
```

```
[root@master ~]# tee /etc/yum.repos.d/kubernetes.repo<<EOF
> [kubernetes]
> name=Kubernetes
> baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
> enabled=1
> gpgcheck=1
> repo_gpgcheck=1
> gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://packages.cloud.google.com/yum/doc/
/rpm-package-key.gpg
> EOF
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://packages.cloud.google.com/yum/doc/r
pm-package-key.gpg
[root@master ~]# █
```

install packages on all machine

```
yum clean all && yum -y makecache
```

```
yum -y install epel-release git curl wget kubelet kubeadm kubectl --disableexcludes=kubernetes
```

start kubelet service on all machine

```
systemctl start kubelet
```

enable kubelet service on all machine

```
Systemctl enable kubelet
```

Add CRI-O repo on all machine

```
OS=CentOS_7
```

```
VERSION=1.22
```

```
curl -L -o /etc/yum.repos.d/devel:kubic:libcontainers:stable.repo
```

```
https://download.opensuse.org/repositories/devel:kubic:libcontainers:stable/\$OS/devel:kubic:libcontainers:stable.repo
```

```
curl -L -o /etc/yum.repos.d/devel:kubic:libcontainers:stable:cri-o:$VERSION.repo
```

```
https://download.opensuse.org/repositories/devel:kubic:libcontainers:stable:cri-o:\$VERSION/\$OS/devel:kubic:libcontainers:stable:cri-o:\$VERSION.repo
```

```

[root@node2 ~]# OS=CentOS_7
[root@node2 ~]# VERSION=1.22
[root@node2 ~]# curl -L -o /etc/yum.repos.d/devel:kubic:libcontainers:stable.repo https://download.opensuse.org/repositories/devel:kubic:libcontainers:stable/$OS/devel:kubic:libcontainers:stable.repo
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left     Speed
100 359 100 359 0 0 188 0 0:00:01 0:00:01 --:--:-- 188
[root@node2 ~]# curl -L -o /etc/yum.repos.d/devel:kubic:libcontainers:stable:cri-o:$VERSION.repo https://download.opensuse.org/repositories/devel:kubic:libcontainers:stable:cri-o:$VERSION/$OS/devel:kubic:libcontainers:stable:cri-o:$VERSION.repo
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left     Speed
100 425 100 425 0 0 694 0 --:--:-- --:--:-- --:--:-- 695
100 426 100 426 0 0 574 0 --:--:-- --:--:-- --:--:-- 0
100 427 100 427 0 0 492 0 --:--:-- --:--:-- --:--:-- 492
100 428 100 428 0 0 431 0 --:--:-- --:--:-- --:--:-- 431
100 429 100 429 0 0 379 0 0:00:01 0:00:01 --:--:-- 379
100 381 100 381 0 0 112 0 0:00:03 0:00:03 --:--:-- 301
[root@node2 ~]# █

```

#Install CRI-O on all machine

```

yum remove docker-ce docker-ce-cli containerd.io
yum install cri-o -y

```

Start and enable Service on all machine

```

systemctl daemon-reload
systemctl start cri-o
systemctl enable cri-o

```

Only on master:-

check status of br_netfilter

```

lsmod | grep br_netfilter

```

```

[root@master ~]# lsmod | grep br_netfilter
br_netfilter          22256  0
bridge                155432  1 br_netfilter

```

Pull image of kubeadm

```

kubeadm config images pull

```

```

[root@master ~]# kubeadm config images pull
[config/images] Pulled registry.k8s.io/kube-apiserver:v1.26.0
[config/images] Pulled registry.k8s.io/kube-controller-manager:v1.26.0
[config/images] Pulled registry.k8s.io/kube-scheduler:v1.26.0
[config/images] Pulled registry.k8s.io/kube-proxy:v1.26.0
[config/images] Pulled registry.k8s.io/pause:3.9
[config/images] Pulled registry.k8s.io/etcd:3.5.6-0
[config/images] Pulled registry.k8s.io/coredns/coredns:v1.9.3

```

create cluster of kubernete

```

kubeadm init --pod-network-cidr=10.85.0.0/16 --upload-certs --control-plane-endpoint=master

```

outpt:-

```
Please note that the certificate-key gives access to cluster sensitive data, keep it secret!
As a safeguard, uploaded-certs will be deleted in two hours; If necessary, you can use
"kubeadm init phase upload-certs --upload-certs" to reload certs afterward.

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join master:6443 --token ifs1ly.62vcct4kwzxhon9c \
--discovery-token-ca-cert-hash sha256:b6c4ba907e6d3a580c30678acc602ed010f0f92e5d9b22134243efe2c6
```

Copy token (last 2nd line)

create a file and paste token

nano join-token

```
GNU nano 2.3.1 File: join-token

kubeadm join master:6443 --token ifs1ly.62vcct4kwzxhon9c \
--discovery-token-ca-cert-hash sha256:b6c4ba907e6d3a580c30678acc602ed010f0f92e5d9b22134243efe2c6
```

Give permission to this file and check

chmod u+x join-token

```
[root@master ~]# chmod u+x join-token
[root@master ~]# ls
anaconda-ks.cfg  Desktop  Downloads  initial-setup-ks.cfg  Music  Public  Videos
join-token      Pictures  Templates
```

copy the file on all node

```
[root@master ~]# rsync join-token root@node1:/root/
[root@master ~]# rsync join-token root@node2:/root/
```

Run token on both node

./join-token

check node on master node

kubectl get nodes

```
[root@master ~]# kubectl get nodes
NAME      STATUS    ROLES    AGE   VERSION
master    Ready     control-plane  34m   v1.26.0
node1     Ready     <none>    17s   v1.26.0
node2     Ready     <none>    9s    v1.26.0
```

Install network plugin:-

Install Calico:

Install the Tigera Calico operator and custom resource definitions.

kubectl create -f

<https://raw.githubusercontent.com/projectcalico/calico/v3.24.5/manifests/tigera-operator.yaml>

Install Calico by creating the necessary custom resource

kubectl create -f

<https://raw.githubusercontent.com/projectcalico/calico/v3.24.5/manifests/custom-resources.yaml>

show all podes

kubectl get pods --all-namespaces

```
[root@master ~]# kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-787d4945fb-dx5zn	1/1	Running	0	43m
kube-system	coredns-787d4945fb-p7657	1/1	Running	0	43m
kube-system	etcd-master	1/1	Running	0	43m
kube-system	kube-apiserver-master	1/1	Running	0	43m
kube-system	kube-controller-manager-master	1/1	Running	0	43m
kube-system	kube-proxy-9sb2x	1/1	Running	0	43m
kube-system	kube-proxy-mxdkg	1/1	Running	0	8m53s
kube-system	kube-proxy-qmt7n	1/1	Running	0	8m45s
kube-system	kube-scheduler-master	1/1	Running	0	43m
tigera-operator	tigera-operator-7795f5d79b-pvvl2	1/1	Running	0	3m34s

show podes in detail information

kubectl get pods --all-namespaces -o wide

```
[root@master ~]# kubectl get pods --all-namespaces -o wide
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP
kube-system	coredns-787d4945fb-dx5zn	1/1	Running	0	44m	10.85.0.3
	master <none>					
kube-system	coredns-787d4945fb-p7657	1/1	Running	0	44m	10.85.0.2
	master <none>					
kube-system	etcd-master	1/1	Running	0	44m	192.168.174.168
	master <none>					
kube-system	kube-apiserver-master	1/1	Running	0	44m	192.168.174.168
	master <none>					
kube-system	kube-controller-manager-master	1/1	Running	0	44m	192.168.174.168
	master <none>					
kube-system	kube-proxy-9sb2x	1/1	Running	0	44m	192.168.174.168
	master <none>					
kube-system	kube-proxy-mxdkg	1/1	Running	0	10m	192.168.174.167
	node1 <none>					
kube-system	kube-proxy-qmt7n	1/1	Running	0	10m	192.168.174.166
	node2 <none>					
kube-system	kube-scheduler-master	1/1	Running	0	44m	192.168.174.168
	master <none>					
tigera-operator	tigera-operator-7795f5d79b-pvvl2	1/1	Running	0	5m5s	192.168.174.166
	node2 <none>					

\$\$\$\$\$ BASIC COMMANDS \$\$\$\$\$\$

for create container

```
kubectl create deployment nginx-image --image=nginx --replicas=1  
replicas =>how many pod want to create
```

create container with port

```
kubectl expose deployment nginx-image --port=9000 --target-port=80 --name=web1  
Target port =>port of container  
Port =>port of pod  
By default port is cluster ip
```

create container with port and type

```
kubectl expose deployment nginx-image --port=9000 --target-port=80 --name=web2  
--type=NodePort
```

for delete the services

```
Kubectl delete services 'service_name'
```

For delete deployment

```
Kubectl delet deployment nginx-image
```

For show pods

```
Kubectl get pods
```

Types of services

```
=> cluster ip  
=> Node port  
=> Load balancer
```

```
# * If we want in internal changes the we want need to make changes in '.yaml' file
```


Worker node(s)

Protocol	Direction	Port Range
TCP	Inbound	10250
TCP	Inbound	30000-32767